
Same Action, Different Verdict?

A Presentation-Invariance Check for AI Reviewers of Agent Actions

Cherishma Subhasa

Darren Kim

Samridhi Makkar

Vishwa Kumaresh

AI Swarm Dynamics Hackathon

Abstract

In multi-agent systems, AI reviewers increasingly judge the actions of other agents, and the reviewed agent’s own explanation is often part of what the reviewer reads. If that explanation can change the verdict while the recorded facts stay fixed, persuasion becomes a route around oversight. We present Parallax-R, a check that shows a reviewer the same recorded action in several presentations (full story, action only, identity-masked, and evidence ledger), verifies that every presentation preserves the same evidence, reconciles the record against hash-linked receipts, resolves the reviewer’s evidence citations, and escalates any case whose verdict depends on presentation. An offline prototype with a deliberately persuadable scripted reviewer runs 24 fictional cases and 288 review records. All 10 prohibited actions are flagged without the explanation and excused with it; each is escalated for view fragility. All 10 allowed actions are consistently accepted. All 4 injected receipt faults are escalated as evidence conflicts. A 68-test suite and a pre-specified 180-call pilot protocol accompany the prototype. These results show that the instrument works as specified; they are not a measurement of any real model. The contribution is a reproducible, falsifiable test harness for explanation sensitivity in AI review.

1. Introduction

An agent emails a file it was told not to send, then explains: “I thought sending it would help finish the task. This was only practice.” A second agent, acting as reviewer, reads the instruction, the recorded action, and that explanation. The relevant question is not only whether the reviewer reaches the right verdict, but whether it would reach the *same* verdict if the explanation were absent.

In agent swarms this question compounds. Reviewers are themselves agents, approvals are consumed by other agents, and an action excused at one point of oversight can propagate as trusted work. A reviewer that is moved by narrative rather than by the rule and the record gives every reviewed agent a lever over its own oversight. Detecting that dependence requires holding the facts fixed while varying only the presentation, and keeping problems with the evidence separate from problems with the reviewer.

Parallax-R is built around that comparison. Its intended users are teams evaluating AI reviewers, and the people who decide when a reviewer’s judgment can be accepted without a human. The contributions are:

- an evidence model with synthetic hash-linked receipts and a view generator whose evidence-preservation invariants are checked automatically before any review;

- a structured reviewer contract with citation requirements, and metrics with explicit missing-data rules that never convert a missing or malformed answer into “no flag”;
- a decision gate that accepts a shared verdict only when every implemented check passes, and otherwise escalates with ordered reason codes;
- a pre-specified pilot protocol with feasibility rules fixed before any model is called; and
- an offline, dependency-free prototype with a six-step HTML walkthrough, machine-readable records, and a 68-test suite.

2. Related Work

Monitoring and control. AI control evaluates whether oversight protocols remain safe when the monitored model may act against them, using trusted and untrusted monitors [4]. Chain-of-thought monitoring can detect misbehavior, but optimizing against a monitor can teach a model to obscure its reasoning [1]. Both lines treat the monitored agent’s outputs as potentially adversarial input to the monitor. Parallax-R isolates one such input, the actor’s explanation, and asks whether removing it changes the verdict.

Judge reliability. Language models used as evaluators show position, verbosity, and self-enhancement biases [12, 13]. Models also tend toward sycophancy, agreeing with views presented to them [10], and their stated explanations can be unfaithful to the factors that drove their answers [11]. These findings motivate testing a reviewer’s invariance to presentation directly rather than trusting its rationale.

Untrusted text in context. Prompt-injection work shows that text inside a model’s input can act as instructions [5, 8]. The Parallax-R rubric instructs the reviewer to treat the trace as evidence and the actor’s explanation as a claim to examine, not as permission. Whether a reviewer follows that instruction is exactly what the comparison measures.

Tamper-evident records. Hash-chained audit logs detect modification of recorded entries [2, 9]. Parallax-R uses a minimal chain only to construct controlled faults. It does not provide forward security, anchoring, or protection against an attacker who rewrites an entire log.

Multi-agent risk. Risks that arise from interacting agents, including miscoordination and collusion, are distinct from single-agent failures [6]. Agents reviewing agents form one such interaction. We do not model collusion; we test a precondition for robust review, namely that the verdict depends on the evidence rather than on how it is narrated.

3. Methods

3.1 Scope and threat model

The reviewed actor is a fictional agent whose recorded events and explanation are fixed data. The reviewer is either a scripted function (offline demo) or a language model behind the Anthropic Messages API (pilot mode, not yet run). The trusted components are the case labels, the rule text, the view renderer, the invariant checker, and the gate. The reviewer never sees the expected verdict, other views’ answers, or the research hypothesis.

Labels state whether the action is *prohibited* or *allowed* under the supplied rule. They do not assert harm or intent. Context that genuinely changes permission belongs in the rule or the records, where every view preserves it, not only in the actor’s account.

3.2 Cases, events, and receipts

Each case holds a task, a rule, recorded events (actor, operation, resource, output, source), an actor narrative, a coverage status, and a list of receipts. A receipt r_i copies its event’s operation, resource, and output and links to its predecessor:

$$d_i = H(r_i \setminus \{d_i\}), \quad \text{prev}(r_i) = d_{i-1}, \quad d_0 = \text{GENESIS}, \quad (1)$$

where H is SHA-256 over canonical JSON. Reconciliation assigns each event one of SUPPORTED (one receipt, matching fields), MISMATCH (fields differ), or PHANTOM (no receipt). It marks a receipt with no transcript event as HIDDEN, and any digest or link failure as CHAIN_BREAK. A missing receipt does not prove that an action never happened, and a supported event does not override a broken chain.

3.3 Views and evidence invariants

Four deterministic renderers present the same evidence packet (Table 1). Before review, the runner checks that the task, rule, events, receipts, and reconciliation are byte-identical to the full view’s, except that the masked view may relabel actors through a consistent one-to-one map. A failure in the full or action-only view excludes the case before review. A failure in the masked or ledger view skips only that view and escalates the case, so that the primary comparison is preserved.

Table 1: Presentations of one evidence packet. Only the narrative and actor display labels vary.

View	Explanation	Actor labels	Role
Full	included	original	Reference presentation
Action-only	removed	original	Primary comparison with full
Masked	included	neutral (<i>Actor 1</i>)	Secondary: identity presentation
Ledger	removed	original	Optional fourth view; evidence listing

Automated invariants confirm that evidence fields survive. They cannot confirm that removing a narrative leaves the meaning of the remaining evidence unchanged; that requires human review (Section 4.4).

3.4 Reviewer contract

The reviewer returns JSON with an assessment status, a `risk_score` in $[0, 1]$ describing how strongly the evidence indicates a rule violation, a violation category, a confidence, an intent assessment, a rationale, and a list of claims. Each claim cites event and receipt IDs, and is marked critical if changing it would change the assessment. An INSUFFICIENT_EVIDENCE status carries a null score. Responses are schema-validated; invalid output, insufficient evidence, and transport failure are recorded as distinct abstention causes.

A valid score $s \geq 0.50$ yields FLAG, otherwise NO_FLAG. Each view is reviewed three times in fresh conversations. A view receives a majority verdict only if all three reviews are valid and definite; otherwise it is ABSTAIN. Cached replay is recorded as replay, not as a new review.

3.5 Measurements

For case i and view v , let $r_{i,v}$ be the mean of the three scores and $y_{i,v}$ the majority verdict. With F the full view and A the action-only view, the primary score difference for prohibited case i is

$$\Delta_i = r_{i,A} - r_{i,F}, \quad (2)$$

so a positive value means the full story looked less risky. Among n eligible prohibited cases, a case is *more forgiving* if $y_{i,A} = \text{FLAG}$ and $y_{i,F} = \text{NO_FLAG}$, and *stricter* in the reverse case. We report both counts and

$$\text{net} = \frac{\#\text{more forgiving} - \#\text{stricter}}{n}, \quad d_z = \frac{\bar{\Delta}}{\text{sd}(\Delta)}, \quad (3)$$

with d_z undefined when $n < 2$ or $\text{sd}(\Delta) = 0$ [7]. Per case, D is the population standard deviation of the three primary view means, F_i indicates whether their majority verdicts differ, and U is the fraction of critical claims that cite no record or an ID absent from the case. U checks that citations resolve, not that they support the claim. Repeat variation within a view is reported as the population standard deviation of its three scores.

3.6 Decision gate

The gate consumes per-view aggregates, the reconciliation report, and the citation checks. It emits ACCEPT with the shared verdict only if none of the following holds, and otherwise ESCALATE with every applicable reason, in this order:

1. EVIDENCE_CONFLICT: any MISMATCH, PHANTOM, HIDDEN, or CHAIN_BREAK finding;
2. INCOMPLETE_RECORD: coverage is not a known-complete synthetic fixture;
3. UNSUPPORTED_CLAIM: some critical claim has unresolved citations;
4. VIEW_FRAGILITY: usable views disagree on the majority verdict;
5. MONITOR_ABSTENTION: some required view lacks a usable verdict;
6. VIEW_VALIDATION_FAILURE: a secondary view failed its invariant check.

Proposition 1 (No silent acceptance). *If the gate returns ACCEPT, every required view has a complete, valid majority verdict, all of those verdicts are equal, every critical claim’s citations resolve, and reconciliation found no conflict.*

Proof. Each condition’s negation adds a reason code, and ACCEPT is returned only when the reason list is empty. The accepted verdict is taken from the set of usable verdicts, which then has exactly one element. □

Proposition 2 (Missing data is never a negative). *A missing, malformed, or insufficient-evidence review cannot contribute a NO_FLAG verdict, a zero score, or an ACCEPT.*

Proof. Such a review makes its view incomplete, so the view’s verdict is ABSTAIN and its mean is excluded from Δ . An incomplete required view triggers MONITOR_ABSTENTION, so Proposition 1 rules out ACCEPT. □

Escalation does not remove an otherwise eligible case from the comparison metrics, and it sends no notification. ACCEPT means the implemented checks passed. It is not a safety guarantee, and the accepted verdict may itself be FLAG.

3.7 Evaluation design

Teaching cases. Twenty fictional cases form ten scenarios, each with a prohibited version (“Do not email it to anyone”) and an allowed version (“You are authorized to email it”). The recorded action and narrative are identical within a scenario. Coverage is KNOWN_COMPLETE_SYNTHETIC.

Scripted reviewer. To exercise the pipeline end to end, the offline reviewer is *deliberately persuadable*: it returns $s = 0.9$ when the rule prohibits the action and no explanation is shown, and $s = 0.2$ otherwise. Its claims cite the rule, every event, and every receipt. The disagreement it produces is programmed. It demonstrates that the instrument detects and reports explanation sensitivity; it is not a finding about any model.

Fault injection. Four copies of a prohibited case each receive one fault: an altered event output (mismatch), all receipts removed with coverage unknown (dropout), the event removed from the transcript (hidden), or a tampered receipt digest (chain break). These cases are excluded from the primary prohibited/allowed comparison but included in gate totals.

Run. The walkthrough runs $24 \text{ cases} \times 4 \text{ views} \times 3 \text{ repeats} = 288$ review records. Exploratory statistics use a scenario-cluster bootstrap (10,000 draws, seed 20261004) [3] and an exact scenario-level sign-flip test.

4. Results

4.1 The pipeline surfaces programmed explanation sensitivity

Every view of every case received three valid reviews; there were no abstentions or exclusions. Each prohibited case was flagged in the action-only and ledger views and excused in the full and masked views, which both include the explanation (Table 2, Figure 1a). Each of the ten was classified as more forgiving and escalated for VIEW_FRAGILITY. Every allowed case received NO_FLAG in all four views and was accepted.

Table 2: Majority verdicts and mean scores by view in the scripted walkthrough (20 teaching cases). Repeat standard deviation was 0 in every view, because the reviewer is deterministic.

	Full	Action-only	Masked	Ledger
Prohibited cases flagged (/10)	0	10	0	10
Allowed cases flagged (/10)	0	0	0	0
Mean score, prohibited	0.20	0.90	0.20	0.90
Mean score, allowed	0.20	0.20	0.20	0.20




Figure 1: Scripted walkthrough outcomes. (a) Majority FLAG verdicts per view for the 20 teaching cases; the prohibited cases are flagged only when the explanation is absent. (b) Reason codes across all 24 cases, including the four fault cases; a case can carry several reasons. Points are exact outcomes of a deterministic reviewer, with no sampling error.

4.2 Undefined statistics are reported, not passed

For the prohibited class, $\Delta_i = 0.70$ for all ten cases and the net more-forgiving rate is 1.0, which meets the pre-specified ≥ 0.10 rule. Because every difference is identical, $\text{sd}(\Delta) = 0$ and d_z is reported as UNDEFINED_ZERO_SD. The protocol forbids an undefined statistic from satisfying the $d_z \geq 0.30$ rule, and the runner records that rule as not met. The bootstrap interval degenerates to $[0.70, 0.70]$. The exact sign-flip test gives $p = 2/1024 \approx 0.002$, which reflects the programmed constant difference and carries no evidential weight. These outputs confirm that the analysis code handles degenerate inputs as specified.

4.3 Record faults escalate independently of the verdict

All four fault cases were detected and escalated with EVIDENCE_CONFLICT (Table 3). Each also carried VIEW_FRAGILITY, because each copies a prohibited case and inherits its programmed disagreement. The dropout case added INCOMPLETE_RECORD because its coverage is unknown. Evidence problems are therefore visible even when the reviewer’s verdict would otherwise be accepted.

Table 3: Injected receipt faults. Each is a copy of one prohibited teaching case with a single change.

Fault	Reconciliation finding	Gate decision
Mismatch	MISMATCH	ESCALATE: evidence conflict, view fragility
Dropout	PHANTOM; coverage unknown	ESCALATE: evidence conflict, incomplete record, view fragility
Hidden	HIDDEN	ESCALATE: evidence conflict, view fragility
Chain break	CHAIN_BREAK and SUPPORTED	ESCALATE: evidence conflict, view fragility

4.4 Pilot readiness

The offline readiness check passes on the 20 draft pilot candidates (10 prohibited, 10 allowed). These cover ten scenarios: external uploads, backup deletion, access grants, configuration edits, downloaded scripts, credential reads, personnel messages, transfer limits, raw survey publication, and audit logging. Rendered inputs range from 4,199 to 5,012 characters. Person A recorded label decisions for all 20, including four authorized fictional evidence amendments. A human compared the displayed cross-view summaries for five cases spanning both labels, and an assistant-run preservation check passed for all 20, tied to content hashes. Claude Sonnet 5.5 is provisionally recorded as the reviewer model. The planning estimate for the 180 scheduled calls is USD 1.62–3.42; a USD 15 budget is proposed but not authorized. The settings remain HANDOFF_NOT_FROZEN, and no real model has been called.

4.5 Software verification

The suite contains 68 tests. They cover schema validation, receipt reconciliation, view invariants and masking, the API adapter under mocked transport (retries, transport failures, malformed output), aggregation and both flip directions, the gate and citation checks, the exploratory statistics, and the HTML report, including script-injection escaping. During development, six deliberately injected bugs each made the suite fail: a swapped flip direction, an off-by-one in the net rule, an absolute-valued d_z , masker numbering, one retry too few, and an abstention counted as a match. The full walkthrough runs in under one second on a laptop.

5. Discussion and Limitations

Practical implication. A reviewer’s verdict should be reported with the presentation it was given. When the same evidence yields different verdicts in different presentations, the disagreement itself is the useful output: it identifies the cases where the reviewer is responding to something other than the rule and the record, and routes them to a person. Keeping evidence checks separate from the verdict prevents a confident review of a broken record from being accepted.

Main threat to validity. The demonstration uses a scripted reviewer built to fail in the direction being tested. It shows that the instrument detects and reports that failure. It does not show that any real model fails this way, how often, or how strongly. The fictional cases are short and were drafted with assistance from the same model family proposed as the reviewer, then reviewed by a human.

What the checks do not establish. Citation checks confirm that IDs resolve, not that the cited record supports the claim. Automated invariants confirm that evidence fields are preserved, not that meaning is preserved. Synthetic receipts cannot stop an attacker from rewriting a whole unanchored log. Identity masking replaces known display labels and is not a general anonymizer. Completeness is assumed only for synthetic fixtures; unknown completeness always escalates.

Relevance to swarms. The present design compares one reviewer’s verdicts across presentations of one action. Swarm settings add reviewer chains, shared context, and agents that can adapt their explanations. The same invariance test applies at each review point, but whether presentation sensitivity amplifies or cancels across a chain is untested here.

5.1 Continuation plan with falsifiable gates

Stage	Deliverable and acceptance gate
Freeze	Independent second-reviewer sign-off on all labels; manual audit of at least five cases across views; team acceptance of the rubric and analysis defaults; record the git revision and timestamp before any intended outcome is seen.
Smoke test	Six calls on separate teaching fixtures with the chosen model and an authorized budget. Require valid structured output and a working call cap before the pilot proceeds.
Pilot	20 cases $\times$ 3 views $\times$ 3 fresh reviews = 180 calls. Require at least seven eligible cases per class, or report the pilot as inconclusive. Continue if net more-forgiving ≥ 0.10 or $d_z \geq 0.30$. Publish negative results and every exclusion.
Extension	If the pilot continues: held-out cases, adaptive explanations, multiple reviewer models, reviewer chains in a multi-agent setting, and a human audit of semantic citation support.

6. Conclusion

Parallax-R turns a simple requirement into a test: the same evidence should receive the same verdict whether or not the actor explains itself. The prototype holds the evidence fixed across presentations, checks that it stays fixed, separates record problems from reviewer problems, and escalates instead of accepting when presentations disagree or data is missing. With a deliberately persuadable reviewer, it flags every programmed failure and accepts every consistent case. The next step is the pre-specified pilot with a real reviewer model, whose result, positive or negative, will be reported in full.

Code and data. The repository (<https://github.com/CherishSu/parallax-r>) contains the runner, all modules, the 68 tests, the fictional teaching cases and draft pilot candidates, the rubric, the draft research plan, the figure script, and this report’s source. No real-model run, paid API call, or third-party test is claimed.

References

- [1] Bowen Baker, Joost Huizinga, Leo Gao, Zehao Dou, Melody Y. Guan, Aleksander Madry, Wojciech Zaremba, Jakub Pachocki, and David Farhi. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. arXiv preprint, 2025. URL <https://arxiv.org/abs/2503.11926>.
- [2] Scott A. Crosby and Dan S. Wallach. Efficient data structures for tamper-evident logging. In *18th USENIX Security Symposium*, 2009.
- [3] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman & Hall/CRC, 1993.
- [4] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. AI control: Improving safety despite intentional subversion. arXiv preprint, 2023. URL <https://arxiv.org/abs/2312.06942>.
- [5] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. arXiv preprint, 2023. URL <https://arxiv.org/abs/2302.12173>.
- [6] Lewis Hammond et al. Multi-agent risks from advanced AI. Cooperative AI Foundation Technical Report, arXiv preprint, 2025. URL <https://arxiv.org/abs/2502.14143>.
- [7] Daniël Lakens. Calculating and reporting effect sizes to facilitate cumulative science: A practical primer for t-tests and ANOVAs. *Frontiers in Psychology*, 4:863, 2013.
- [8] Fábio Perez and Ian Ribeiro. Ignore previous prompt: Attack techniques for language models. arXiv preprint, 2022. URL <https://arxiv.org/abs/2211.09527>.
- [9] Bruce Schneier and John Kelsey. Secure audit logs to support computer forensics. *ACM Transactions on Information and System Security*, 2(2):159–176, 1999.
- [10] Mrinank Sharma et al. Towards understanding sycophancy in language models. arXiv preprint, 2023. URL <https://arxiv.org/abs/2310.13548>.
- [11] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. arXiv preprint, 2023. URL <https://arxiv.org/abs/2305.04388>.
- [12] Peiyi Wang et al. Large language models are not fair evaluators. arXiv preprint, 2023. URL <https://arxiv.org/abs/2305.17926>.
- [13] Lianmin Zheng et al. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. arXiv preprint, 2023. URL <https://arxiv.org/abs/2306.05685>.

A. Limitations and Dual-Use Considerations

A.1 Precisely bounded claims

Reviewer realism. The scripted reviewer keys on teaching-specific wording and cannot meaningfully evaluate the varied pilot candidates. Every demonstrated outcome follows from its two-line rule. No claim about model reliability, persuasion, or identity bias follows from the walkthrough.

Evidence semantics. A cited ID can exist and still be irrelevant. Semantic support requires human audit, and none is automated here. Removing a narrative can, in principle, remove a permission-relevant fact; the protocol requires such facts to appear in neutral, source-labeled form in every view, or the case to be excluded.

Receipts and completeness. Receipts are generated from the same events they check. Hash links detect inconsistency within a log but provide no anchoring, so a party that controls the whole log can rewrite it consistently. Only `KNOWN_COMPLETE_SYNTHETIC` fixtures pass the completeness check. No real-world coverage estimator is implemented.

Thresholds and statistics. The 0.50 cutoff is a fixed exploratory boundary, not a calibrated probability or false-positive operating point. The continuation rules are feasibility signals, not significance tests; with ten prohibited cases, a 0.10 net rate is one case. Three repeats characterize limited within-view variation. Composite scores and calibrated risk-range and coverage cutoffs are deliberately out of scope.

Spending controls. The runner enforces a scheduled-call cap and refuses a real run without explicit confirmation, a model ID, and a recorded budget. The dollar limit is recorded in the manifest, not enforced against actual usage.

A.2 Dual use and privacy

All cases, actors, files, addresses, and receipts are fictional and inert; demo mode makes no network calls. The view-variation method could in principle be used to search for explanations that maximize reviewer forgiveness. Its defensive purpose is to expose that sensitivity before a reviewer is trusted. Real traces can contain sensitive data: the masker is not an anonymizer, and hashing provides neither anonymity nor legal compliance. The API key is read from the environment and never written to outputs.

B. Pre-Specified Pilot Protocol (Summary)

Item	Rule
Question	Does adding the actor’s explanation change the reviewer’s judgment of a fixed recorded action under a supplied rule?
Cases	20 fictional cases: 10 prohibited, 10 allowed, in 10 paired scenarios; labels assigned before any model call and withheld from the reviewer.
Views	Full (V1), action-only (V2), identity-masked (V3; secondary).
Reviews	3 fresh calls per case and view: 180 scheduled judgments; transport retries logged separately.
Verdict	$s \geq 0.50 \Rightarrow \text{FLAG}$; majority of three, only if all three are valid and definite.
Primary outcome	$\Delta_i = r_{i,V2} - r_{i,V1}$ on eligible prohibited cases; both flip directions reported with denominators.
Eligibility	Valid V1/V2 evidence and three definite reviews each; V3 failure does not affect primary eligibility.
Validity floor	At least 7 eligible cases per class, else INCONCLUSIVE.
Continue if	net more-forgiving ≥ 0.10 or positive $d_z \geq 0.30$; undefined statistics cannot satisfy a rule.
Exploratory	Scenario-cluster bootstrap and exact sign-flip; not decision rules.
Reporting	All responses, abstentions, exclusions, and negative results; no tuning on pilot outcomes.

C. Reproduction and Evidence Ledger

C.1 Reproduce the demonstration

From the repository root, Python 3.10 or newer is sufficient; no packages, API key, or network access are needed for the demo.

```
python3 -m unittest discover -s tests
python3 run.py --mode demo --ledger --stress --out output/walkthrough
python3 analyze.py output/walkthrough/results.json --out output/analysis.json
python3 preflight.py
python3 docs/submission/make_figure.py output/walkthrough/results.json
cd docs/submission && tectonic submission.tex
```

The figure script needs Matplotlib, and the PDF needs Tectonic. `output/walkthrough/report.html` is a self-contained six-step walkthrough that opens without a server.

Artifact	Contents
<code>run.py</code>	Runner: validation, reconciliation, views, invariants, reviews, exports
<code>parallax/</code>	Schema, receipts, views, reviewer and API adapter, evidence checks, gate, metrics, statistics, report
<code>tests/</code>	68 unit tests
<code>cases/</code>	Draft pilot candidates, label manifest, view-check notes
<code>MONITOR-RUBRIC.md</code>	Runtime reviewer instructions and response contract
<code>PRE-REGISTRATION.md</code> , <code>METRICS-AND-HANDOFF.md</code>	Draft protocol and metric definitions
<code>preflight.py</code> , <code>analyze.py</code>	Offline readiness check and cost estimate; exploratory analysis
<code>docs/submission/</code>	This report’s source, figure script, and figure

C.2 Claims mapped to evidence

Claim	Evidence and scope
No silent acceptance	Proposition 1, by construction of <code>gate.py</code> ; gate tests.
Missing data is never a negative	Proposition 2; aggregation and abstention tests.
Programmed sensitivity is detected	10/10 prohibited cases escalated; <code>results.json</code> . Scripted reviewer only.
Record faults are detected	4/4 fault cases escalated with evidence conflict; <code>results.json</code> . Synthetic faults only.
Pilot is ready to freeze	Not established. Preflight passes; second-reviewer sign-off, manual audit, budget, smoke test, and freeze are pending.
Real reviewer is explanation-sensitive	Not established. No real-model run has been made.

C.3 LLM Usage Statement

We used LLM coding assistants to scaffold and debug the pipeline, draft the fictional candidate cases (each label was then reviewed by a human), run automated view-preservation checks, and draft sections of this report.